Consensus I

FLP Impossibility, Paxos

جامعة الملك عبدالله
للعلوم والتقنية
King Abdullah University of
Science and Technology

CS 240: Computing Systems and Concurrency
## Lecture 9

Marco Canini

Credits: Michael Freedman and Kyle Jamieson developed much of the original material.

Recall our 2PC commit problem

Client C

Transaction Coordinator TC

Bank   A   B

1. **C → TC:** *"go!"*

2. **TC → A, B:** *"prepare!"*

3. **A, B → P:** *"yes"* or *"no"*

4. **TC → A, B:** *"commit!"* or *"abort!"*

Recall our 2PC commit problem

Client C

**Transaction Coordinator TC**

Bank   A   B

- Who acts as TC?

- Which server(s) own the account of A? B?

- Who takes over if TC fails? What about if A or B fail?

Doing failover "correctly" isn't easy

Transaction Coordinator TC

Which node takes over as backup?

Doing failover "correctly" isn't easy

Transaction Coordinator TC — 1

Okay, so specify some ordering

(manually, using some identifier)

2   3

Doing failover "correctly" isn't easy

Transaction Coordinator TC 1

But who determines
if 1 failed?

2   3

Doing failover "correctly" isn't easy

Easy, right?
Just ping and timeout!

**Transaction
Coordinator TC** 1   2   3

Doing failover "correctly" isn't easy

Is the server or the network actually dead/slow?

Transaction Coordinator TC

1

1

2

What can go wrong?

Two nodes think they are TC:

"Split brain" scenario

Transaction Coordinator TC

1   1

What can go wrong?

Two nodes think they are TC:

"Split brain" scenario

Transaction Coordinator TC

1

1

What can go wrong?

Transaction Coordinator TC

Safety invariant:
Only 1 node is TC at any single time

1

Another problem:
A and B need to know (and agree upon) who the TC is…

Consensus

Definition:

1. A general agreement about something

2. An idea or opinion that is shared by all the people in a group

Origin: Latin, from *consentire*

Consensus

Given a set of processors, each with an initial value:

- **Termination:** All non-faulty processes eventually decide on a value

- **Agreement:** All processes that decide do so on the same value

- **Validity:** The value that has been decided must have proposed by some process

Consensus used in systems

Group of servers attempting:

- Make sure all servers in group receive the same updates in the same order as each other

- Maintain own lists (views) on who is a current member of the group, and update lists when somebody leaves/fails

- Elect a leader in group, and inform everybody

- Ensure mutually exclusive (one process at a time only) access to a critical resource like a file

Step one: Define your system model

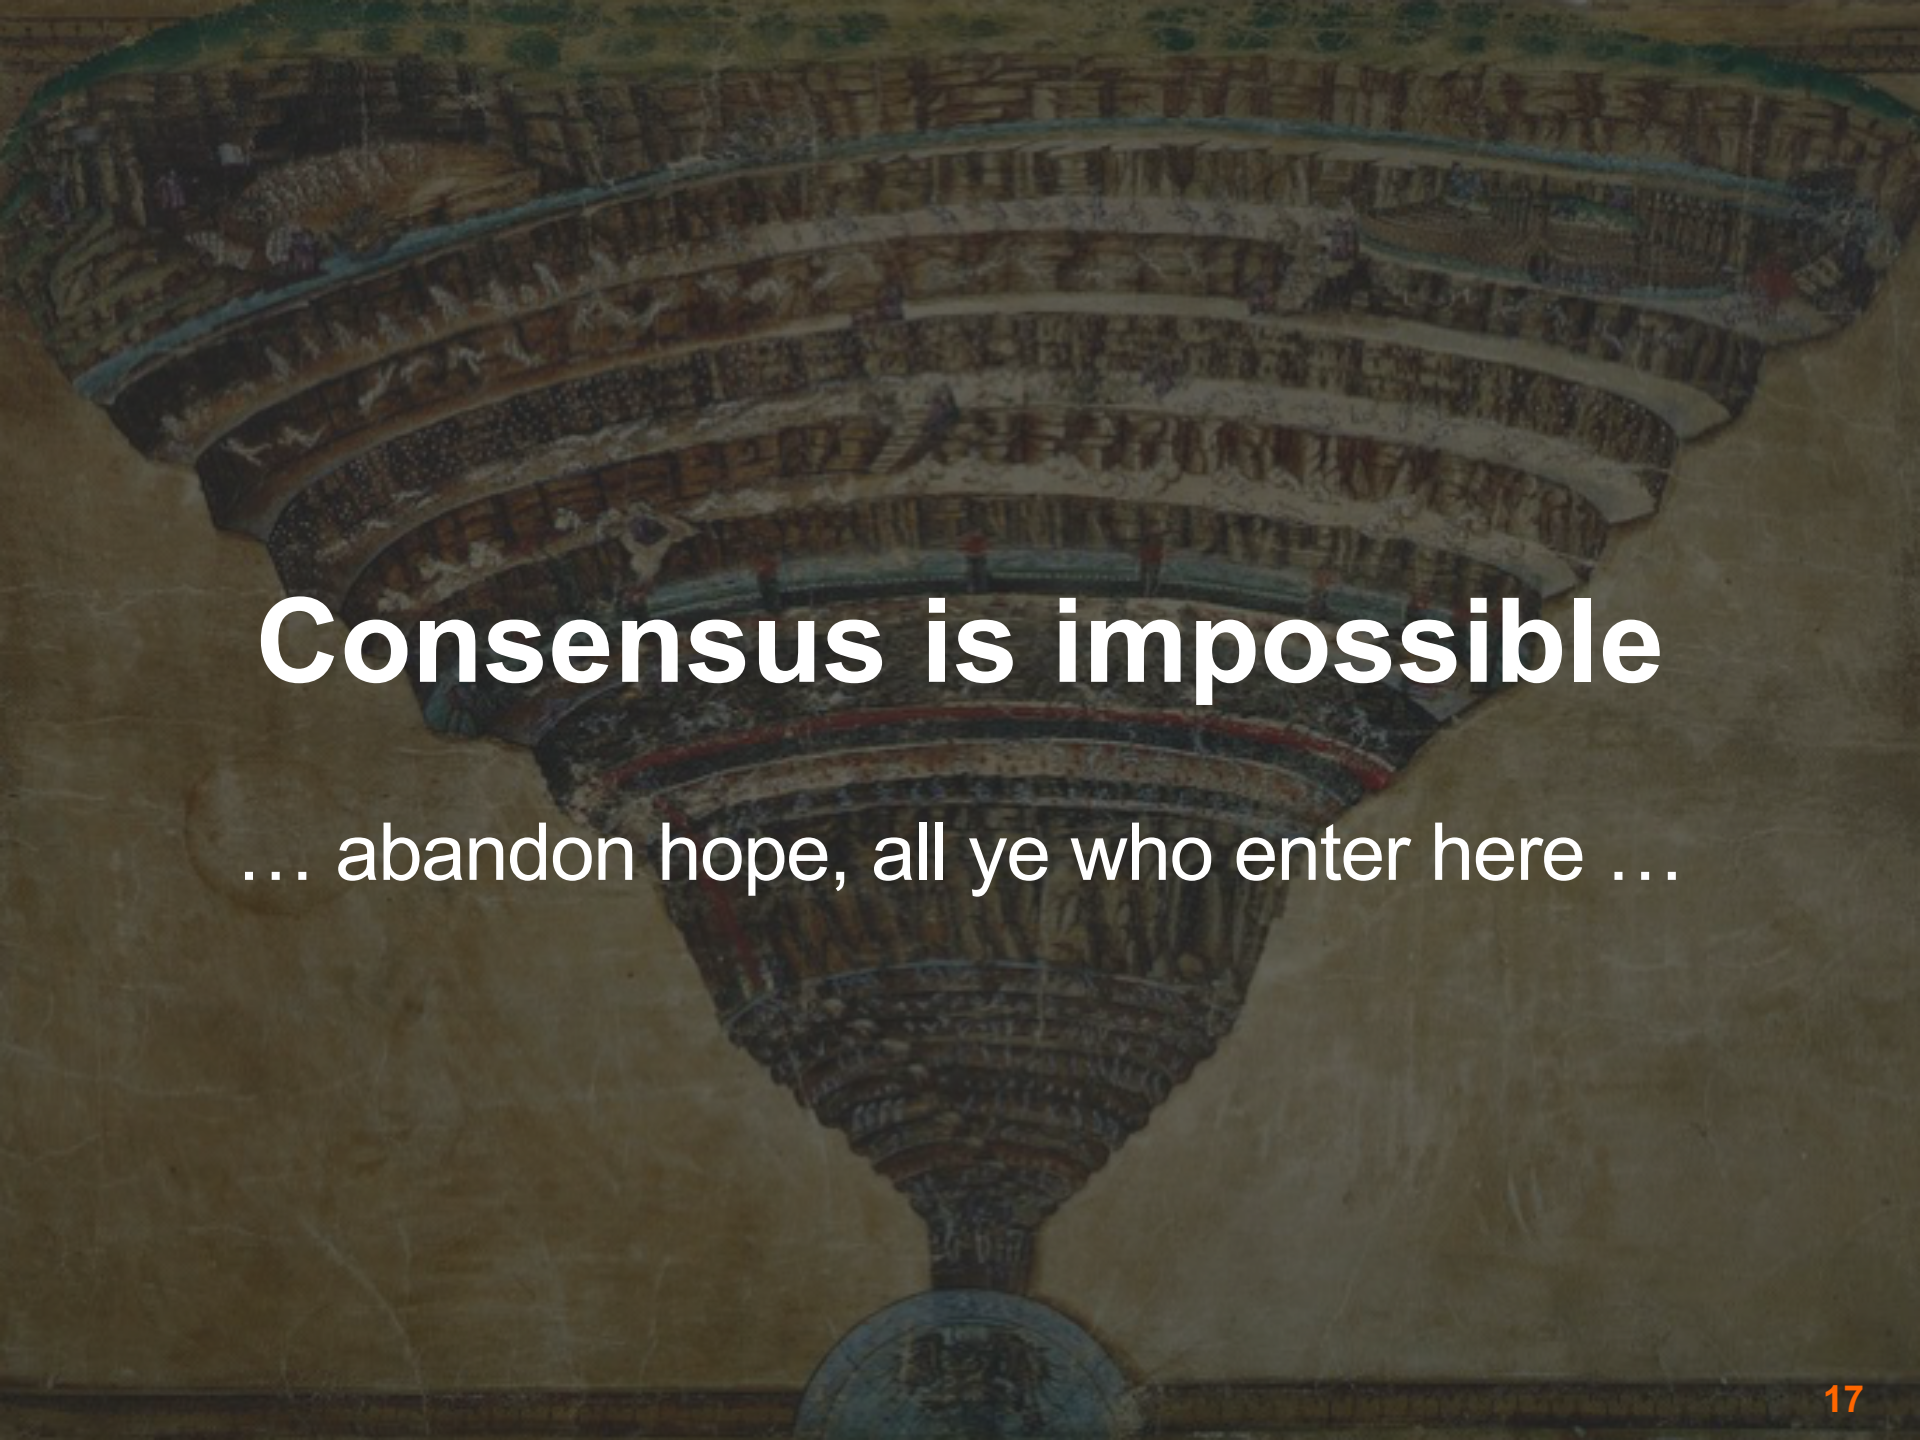- Network model:

  – Synchronous (time-bounded delay) or asynchronous (arbitrary delay)

  – Reliable or unreliable communication

  – Unicast or multicast communication

- Node failures:

  – Fail-stop (correct/dead) or Byzantine (arbitrary)

Step one: Define your system model

- Network model:

  – Synchronous (time-bounded delay) or asynchronous (arbitrary delay)

  – Reliable or unreliable communication

  – Unicast or multicast communication

- Node failures:

  – Fail-stop (correct/dead) or Byzantine (arbitrary)

Consensus is impossible

… abandon hope, all ye who enter here …

"FLP" result

Impossibility of Distributed Consensus with One Faulty Process

MICHAEL J. FISCHER

Yale University, New Haven, Connecticut

NANCY A. LYNCH

Massachusetts Institute of Technology, Cambridge, Massachusetts

AND

MICHAEL S. PATERSON

University of Warwick, Coventry, England

Abstract. The consensus problem involves an asynchronous system of processes, some of which may be unreliable. The problem is for the reliable processes to agree on a binary value. In this paper, it is shown that every protocol for this problem has the possibility of nontermination, even with only one faulty process. By way of contrast, solutions are known for the synchronous case, the "Byzantine Generals" problem.

- No deterministic 1-crash-robust consensus algorithm exists for asynchronous model

- Holds even for "weak" consensus (i.e., only *some* process needs to decide, not *all*)

- Holds even for only two states: 0 and 1

Main technical approach

- Initial state of system can end in decision "0" or "1"

- Consider 5 processes, each in some initial state

$$[1,1,0,1,1] \rightarrow 1$$
$$[1,1,0,1,0] \rightarrow ?$$
$$[1,1,0,0,0] \rightarrow ?$$
$$[1,1,1,0,0] \rightarrow ?$$
$$[1,0,1,0,0] \rightarrow 0$$

Must exist two configurations here which differ in decision

Main technical approach

- Initial state of system can end in decision "0" or "1"

- Consider 5 processes, each in some initial state

[1,1,0,1,1]   →   1
[1,1,0,1,0]   →   1
[1,1,0,0,0]   →   1
[1,1,1,0,0]   →   0
[1,0,1,0,0]   →   0

Assume decision differs between these two processes

Main technical approach

- Goal: Consensus holds in face of 1 failure

**One of these configs must be "bi-valent":
Both futures possible**

$$[1,1 \quad 0,0] \quad \rightarrow \quad 1 \mid 0$$

$$[1,1 \quad 0,0] \quad \rightarrow \quad 0$$

Main technical approach

- Goal: Consensus holds in face of 1 failure

One of these configs must be "bi-valent":
Both futures possible

$$[\, 1,1 \qquad 0,0 \,] \quad \rightarrow \quad 1$$
$$[\, 1,1 \qquad 0,0 \,] \quad \rightarrow \quad 0 \mid 1$$

- Key result: All bi-valent states can remain in bi-valent states after performing some work

You won't believe this one trick!

1.   System thinks process *p* crashes, adapts to it…

2.   But then *p* recovers and *q* crashes…

3.   Needs to wait for *p* to rejoin, because can only handle 1 failure, which takes time for system to adapt …

4.   *… repeat ad infinitum …*

All is not lost…

- But remember
  - "Impossible" in the formal sense, i.e., "there does not exist"
  - Even though such situations are extremely unlikely …

- Circumventing FLP Impossibility
  - Probabilistically
  - Randomization
  - Partial Synchrony (e.g., "failure detectors")

Why should you care?

Werner Vogels, Amazon CTO

Job openings in my group

What kind of things am I looking for in you?

*"**You know your distributed systems theory**: You know about logical time, snapshots, stability, message ordering, but also acid and multi-level transactions. **You have heard about the FLP impossibility argument.** You know why failure detectors can solve it (but you do not have to remember which one diamond-w was). **You have at least once tried to understand Paxos by reading the original paper."**

Paxos

- Safety

  - Only a single value is chosen

  - Only a proposed value can be chosen

  - Only chosen values are learned by processes

- Liveness ***

  - Some proposed value eventually chosen if fewer than half of processes fail

  - If value is chosen, a process eventually learns it

Roles of a Process

- Three conceptual roles

  - Proposers propose values

  - Acceptors accept values, where chosen if majority accept

  - Learners learn the outcome (chosen value)

- In reality, a process can play any/all roles

Strawman

- 3 proposers, 1 acceptor

    – Acceptor accepts first value received

    – No liveness on failure

- 3 proposals, 3 acceptors

    – Accept first value received, acceptors choose common value known by majority

    – But no such majority is guaranteed

Paxos

- Each acceptor accepts *multiple proposals*

  - Hopefully one of multiple accepted proposals will have a majority vote (and we determine that)

  - If not, rinse and repeat (more on this)

- How do we select among multiple proposals?

- Ordering: proposal is tuple (proposal #, value) = (n, v)

  - Proposal # strictly increasing, globally unique

  - Globally unique? Trick: set low-order bits to proposer's ID

Paxos Protocol Overview

- Proposers:

  1. Choose a proposal number n

  2. Ask acceptors if any accepted proposals with $n_a < n$

  3. If existing proposal v_a returned, propose same value (n, v_a)

  4. Otherwise, propose own value (n, v)

  Note altruism: goal is to reach consensus, not "win"

- Accepters try to accept value with highest proposal n

- Learners are passive and wait for the outcome

Paxos Phase 1

- Proposer:
  - Choose proposal number n, send <prepare, n> to acceptors

- Acceptors:
  - If $n > n_h$
    - $n_h = n$ ← promise not to accept any new proposals n' < n
    - If no prior proposal accepted
      - Reply < promise, n, Ø >
    - Else
      - Reply < promise, n, (n_a, v_a) >
  - Else
    - Reply < prepare-failed >

Paxos Phase 2

- Proposer:

  - If receive promise from majority of acceptors,

    - Determine v_a returned with highest n_a, if exists
    - Send <accept, $(n, v_a \,\|\, v)$> to acceptors

- Acceptors:

  - Upon receiving (n, v), if $n \geq n_h$,

    - Accept proposal and notify learner(s)
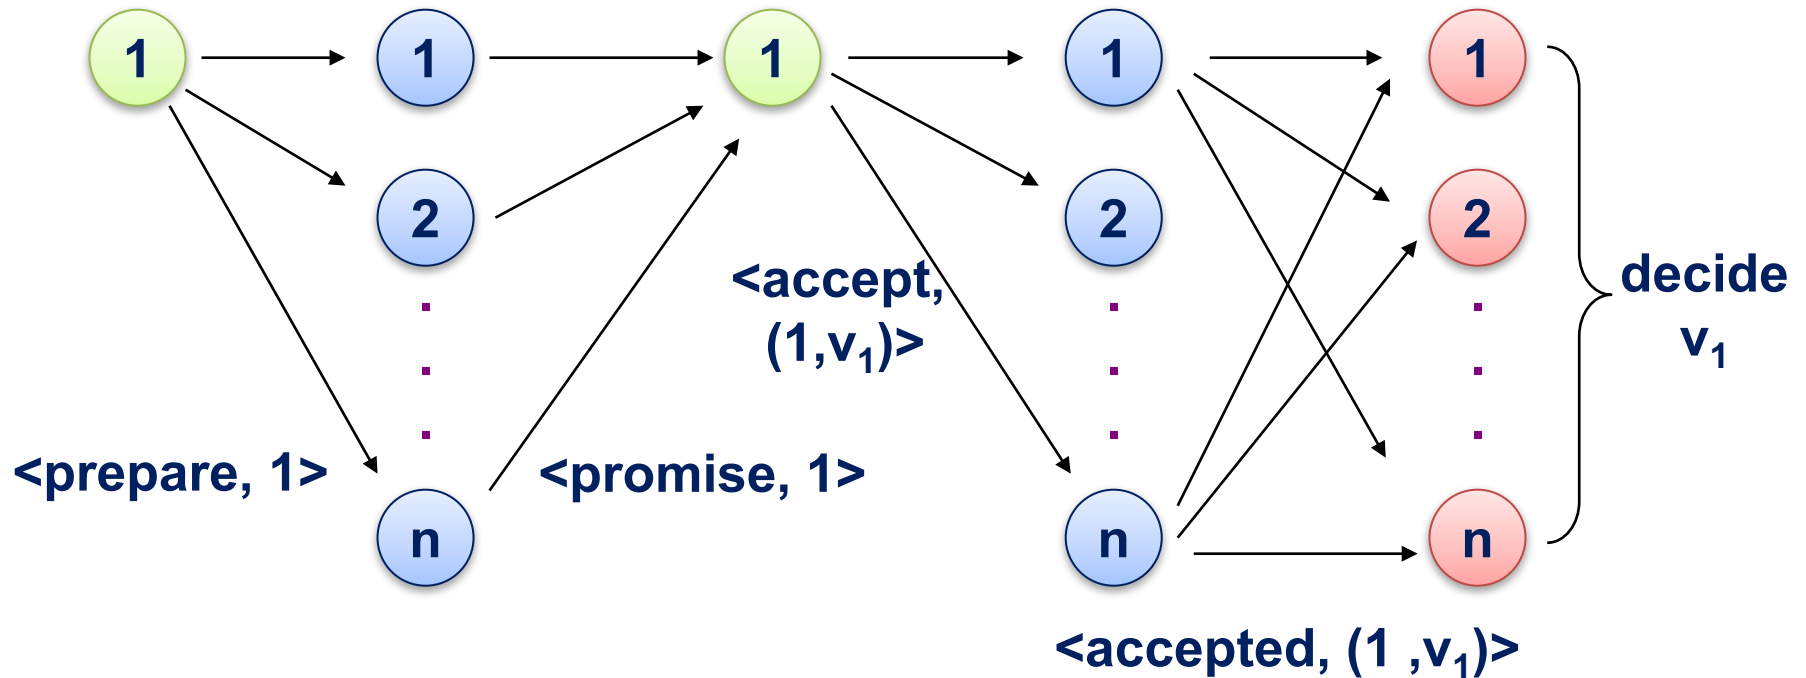
      $n_a = n_h = n$

      $v_a = v$

Paxos Phase 3

- Learners need to know which value chosen

- Approach #1
  - Each acceptor notifies all learners
  - More expensive

- Approach #2
  - Elect a "distinguished learner"
  - Acceptors notify elected learner, which informs others
  - Failure-prone

Paxos: Well-behaved Run


<prepare, 1>

<promise, 1>

<accept, (1,v_1)>

<accepted, (1 ,v_1)>

decide v_1

Paxos is safe
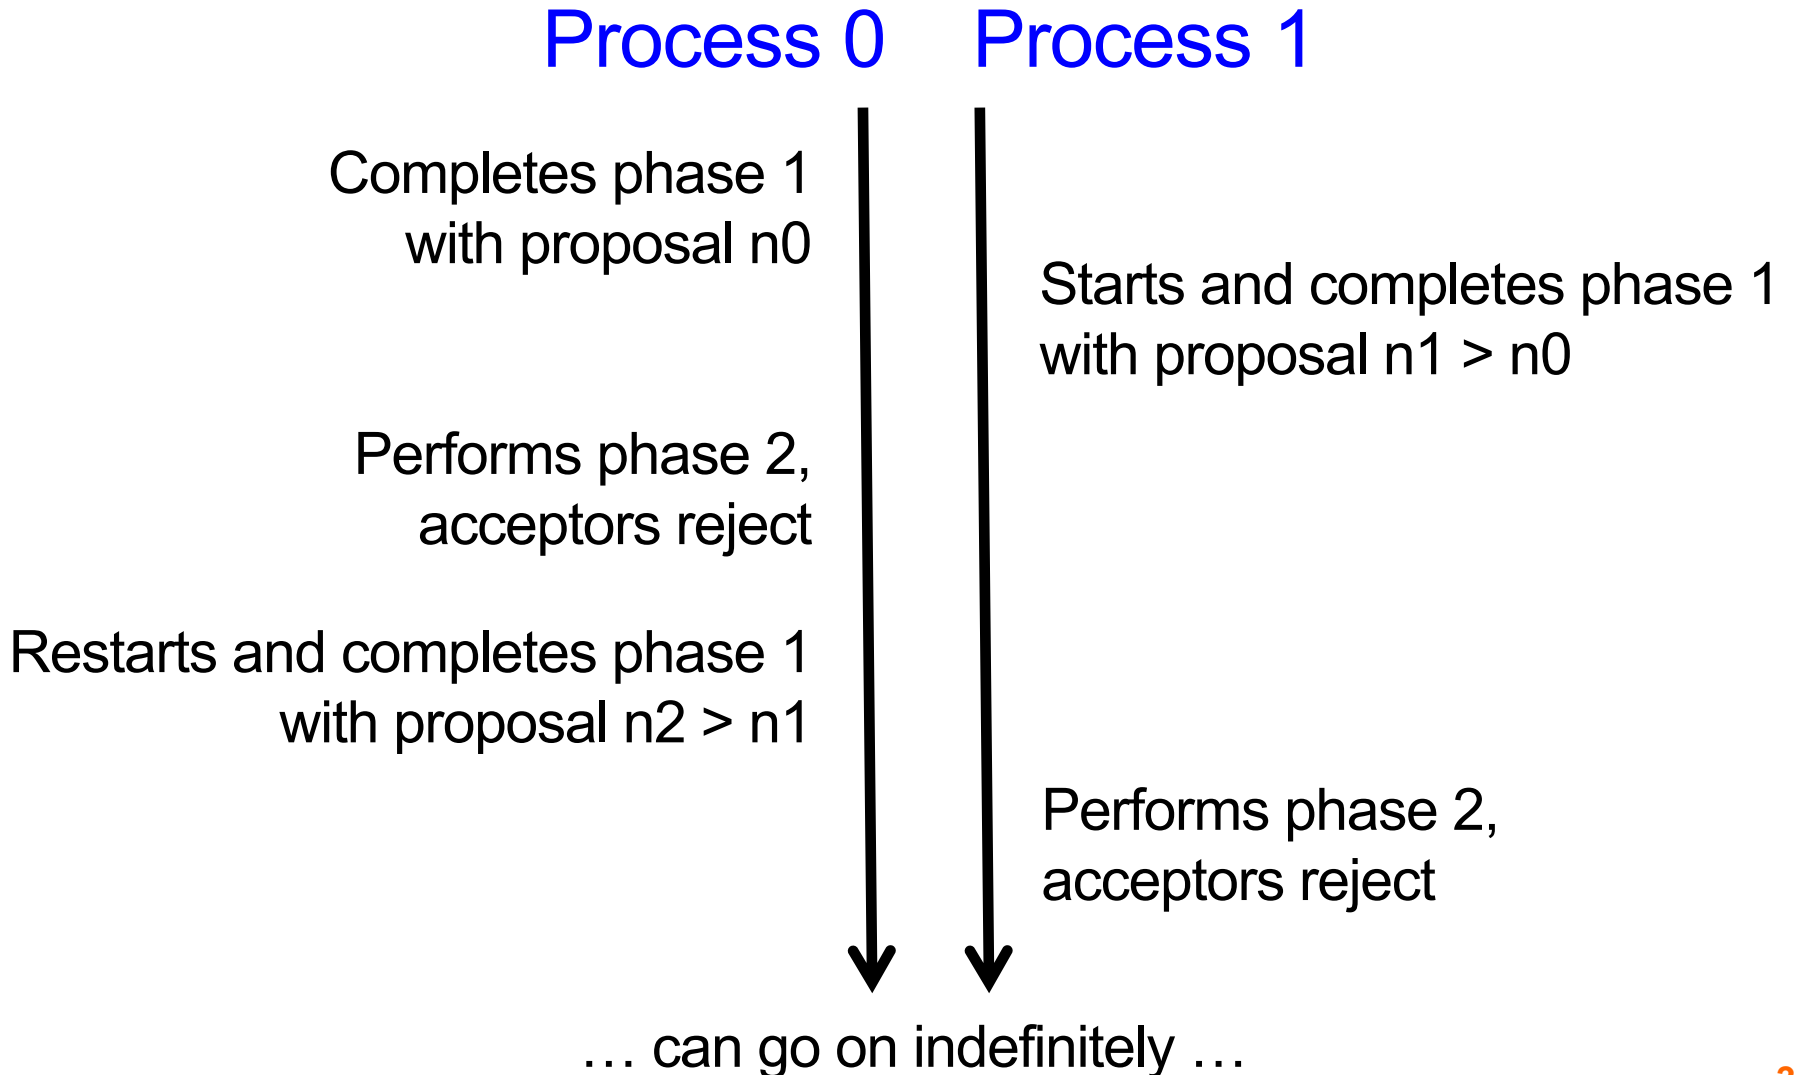
- Intuition: if proposal with value v decided, then every higher-numbered proposal issued by any proposer has value v.

Majority of acceptors accept *(n, v):*

v is decided

Next prepare request with proposal n+1

Race condition leads to liveness problem

Process 0   Process 1

Completes phase 1
with proposal n0

Starts and completes phase 1
with proposal n1 > n0

Performs phase 2,
acceptors reject

Restarts and completes phase 1
with proposal n2 > n1

Performs phase 2,
acceptors reject

… can go on indefinitely …

Paxos with leader election

- Simplify model with each process playing all three roles

- If elected proposer can communicate with a majority, protocol guarantees liveness

- Paxos can tolerate failures $f < N / 2$

Using Paxos in system

Leader election to decide transaction coordinator

L   2   3

Using Paxos in system

New leader election protocol

L   **L**~~new~~   **3**

Still have split-brain scenario!

The Part-Time Parliament

Leslie Lamport

This article appeared in *ACM Transactions on Computer Systems 16*, 2 (May 1998), 133-169. Minor corrections were made on 29 August 2000.

- Tells mythical story of Greek island of Paxos with "legislators" and "current law" passed through parliamentary voting protocol

- Misunderstood paper: submitted 1990, published 1998

- Lamport won the Turing Award in 2013

The Paxos story…

As Paxos prospered, legislators became very busy.

Parliament could no longer handle all details of government, so a bureaucracy was established.

Instead of passing a decree to declare whether each lot of cheese was fit for sale, Parliament passed a decree appointing a cheese inspector to make those decisions.

Cheese inspector ≈ leader
using quorum-based voting protocol

The Paxos story…

Parliament passed a decree making Δῐκστρα the first cheese inspector. After some months, merchants complained that Δῐκστρα was too strict and was rejecting perfectly good cheese.

Parliament then replaced him by passing the decree

  1375: Γωυδα is the new cheese inspector

But Δῐκστρα did not pay close attention to what Parliament did, so he did not learn of this decree right away.

There was a period of confusion in the cheese market when both Δῐκστρα and Γωυδα were inspecting cheese and making conflicting decisions.

Split-brain!

The Paxos story…

To prevent such confusion, the Paxons had to guarantee that a position could be held by at most one bureaucrat at any time.

To do this, a president included as part of each decree the time and date when it was proposed.

A decree making Δἴκστρα the cheese inspector might read

2716: 8:30 15 Jan 72 – Δἴκστρα is cheese inspector for 3 months.

Leader gets a lease!

The Paxos story…

A bureaucrat needed to tell time to determine if he currently held a post. Mechanical clocks were unknown on Paxos, but Paxons could tell time accurately to within 15 minutes by the position of the sun or the stars.

If Δĭκστρα's term began at 8:30, he would not start inspecting cheese until his celestial observations indicated that it was 8:45.

Handle clock skew:

Lease doesn't end until expiry + max skew

Solving Split Brain

New leader election protocol

L L_{new} 3

Solution

If L isn't part of majority electing L_{new}

L_{new} waits until L's lease expires before accepting new ops

Next lecture: Sunday

Other consensus protocols with group membership + leader election at core

- Viewstamped Replication

- RAFT (assignment 3)